

# Data Analysis Assignment 1

Cameron Smith

June 12, 2021

## 1 Question 1

### 1.1 Task 1

We want to determine the decay rate  $\alpha$  for the radioactive material. A plot of the data is shown in Figure 1. To find  $\alpha$ , we use a Stan model with parameters

$$\begin{aligned}\alpha &\sim \mathcal{U}(0, \alpha_{max}) , \\ N_{0,i} &\sim \mathcal{U}(0, N_{0,max})\end{aligned}$$

and

$$t_{0,i} \sim \mathcal{U}(t_{0,min}, t_{0,max}) ,$$

and model

$$N_{1,i} \sim \mathcal{N}(N_{0,i} \exp(-\alpha[t_{i,1} - t_{i,0}]), \sigma_i) ,$$

where  $N_{0,i}$  is the material in the widget  $i$  at the time of manufacture,  $N_{1,i}$  is the material in the widget  $i$  at the time of observation,  $t_{1,i}$ ,  $N_{0,max} = 20\text{ g}$  is the maximum amount of material in a newly manufactured widget,  $t_{0,min}$  and  $t_{0,max}$  refer to the time at the beginning and end of manufacturing respectively,  $\sigma_i$  is the uncertainty in each measurement, and

$$\alpha_{max} = -\frac{1}{\Delta t_{\text{delay}}} \log\left(\frac{\min(N_1)}{\max(N_0)}\right)$$

is the maximum possible decay rate that could produce any of our data.

The model was optimised then sampled over. The sampling chains and resulting log posterior distribution is shown in Figure 2. Through sampling we found a value for  $\alpha$  of

$$\alpha = (5.8 \pm 0.5) \times 10^{-8} \text{ s}^{-1} .$$

### 1.2 Task 2

We now want to consider bias in one (unknown) detector. To account for this, we introduce new parameters  $b$ , the level of bias, and  $\theta_k$ , the probability that detector  $k$  has a bias. We have

$$\theta_k \sim \text{Multinomial}(K)$$

and

$$b \sim \mathcal{U}(-\infty, \infty) .$$

The mean of each observed quantity by a given detector ( $N_{\text{obs},k,i}$ ) is then related to the true quantity ( $N_{\text{true},k,i}$ ) by

$$\mathbb{E}(N_{\text{obs},k,i}) = \mathbb{E}(N_{\text{true},k,i} + B_k),$$

where  $B_k$  is the random variable corresponding to the bias, given by the distribution:

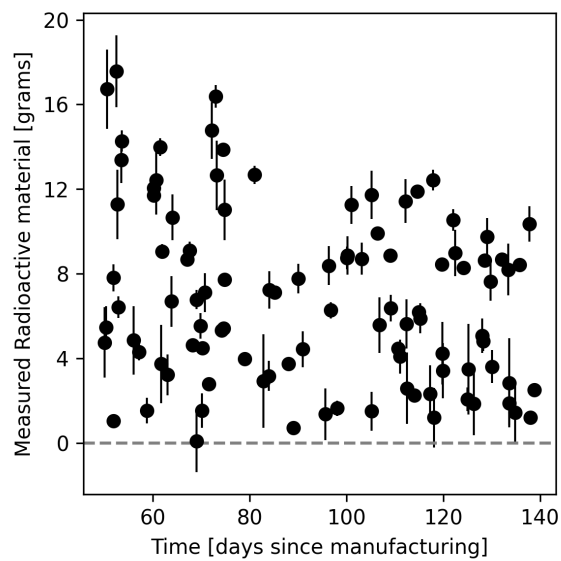


Figure 1: Radioactivity as a function of time

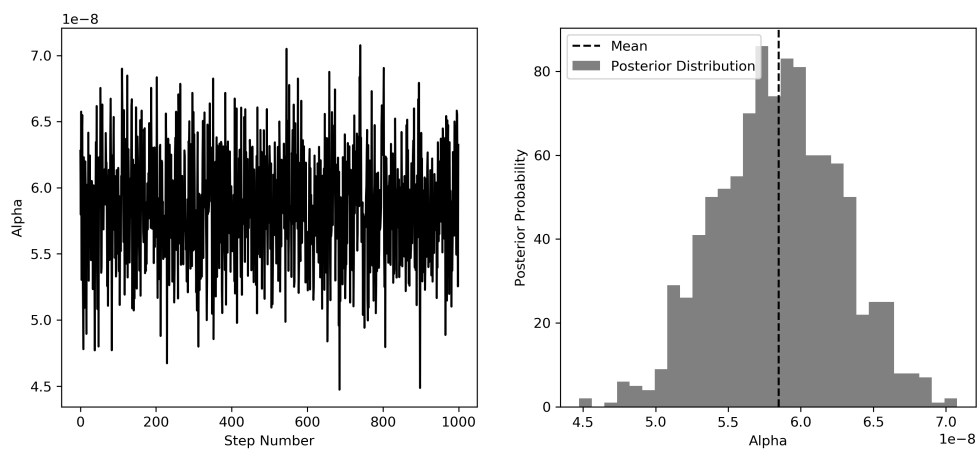


Figure 2: Sampling chains (left) and posterior probability (right) of the decay rate  $\alpha$ .




Figure 3: Sampling chains (left) and posterior probability (right) for the parameters  $\theta_k$ .  $\theta_1$ ,  $\theta_2$  and  $\theta_3$  correspond to detectors A, B and C respectively.

$$\frac{B_k}{P(B_k)} \parallel \begin{array}{c|c} b & 0 \\ \hline \theta_k & 1 - \theta_k \end{array}$$

We can therefore simplify our expression as follows

$$\begin{aligned} \mathbb{E}(N_{\text{obs},k,i}) &= \mathbb{E}(N_{\text{true},k,i}) + \mathbb{E}(B_k), \\ &= N_{0,i} \exp(-\alpha[t_{i,1} - t_{i,0}]) + \theta_k b. \end{aligned}$$

Our model therefore becomes:

$$N_{1,k,i} \sim \mathcal{N}(N_{0,i} \exp(-\alpha[t_{i,1} - t_{i,0}]) + \theta_k b, \sigma_i).$$

Our new model was again implemented in Stan, optimised and sampled over. Figures 3, 4 and 5 show the sampling chains and resulting posterior distributions of parameters  $\theta$ ,  $b$  and  $\alpha$  respectively. From Figure 3 we can see that detector B is likely to be biased. The posterior distribution shown in Figure 4 gives the level of bias as:

$$\alpha = (4.2 \pm 0.8) \text{ g}.$$

Now that we have accounted for the bias, we get a new value for  $\alpha$ :

$$\alpha = (8.2 \pm 0.7) \times 10^{-8} \text{ s}^{-1}.$$

The code for Question 1 is shown in Appendix A.

## 2 Question 2

### 2.1 Task 1

Figure 6(a) shows the brightness of a star as a function of time. Due to stellar rotation, the signal is quasi-periodic. Figure 6(b) shows a Lomb-Scargle periodogram of the data. The peaks with a power greater than 100 are:

Frequency [day <sup>-1</sup> ]	Power	Period [days]
0.123	303	51
0.229	1672	27
0.327	161	19
0.421	922	14
0.605	134	10




Figure 4: Sampling chains (left) and posterior probability (right) for the bias  $b$ .






Figure 5: Sampling chains (left) and posterior probability (right) for the decay rate  $\alpha$ , after accounting for a bias.



(a)



(b)

Figure 6: (a) The brightness of a star as a function of time. Error bars have been omitted for clarity. (b) The Lomb-Scargle periodogram of the data. Peaks with a power greater than 100 are indicated with vertical lines.

where the periods were calculated according to

$$T = \frac{2\pi}{f},$$

where  $f$  is the frequency. The periodogram is dominated by the frequencies at  $0.229 \text{ day}^{-1}$  and  $0.421 \text{ day}^{-1}$ , which possibly correspond to different stellar rotation rates at different latitudes.

## 2.2 Task 2

Now that we know the periodicity of the data, we can fit the data and make predictions of the future brightness using a gaussian process. To do this we used the following combination of kernels

$$k(\mathbf{x}_i, \mathbf{x}_j) = k_1(\mathbf{x}_i, \mathbf{x}_j) (k_2(\mathbf{x}_i, \mathbf{x}_j) + k_3(\mathbf{x}_i, \mathbf{x}_j))$$

where,

$$\begin{aligned} k_1(\mathbf{x}_i, \mathbf{x}_j) &= \exp\left(-\frac{(x - x')^2}{2}\right), \\ k_2(\mathbf{x}_i, \mathbf{x}_j) &= \exp\left(-\Gamma_2 \sin^2\left[\frac{\pi}{P_2} |x_i - x_j|\right]\right), \text{ and} \\ k_3(\mathbf{x}_i, \mathbf{x}_j) &= \exp\left(-\Gamma_3 \sin^2\left[\frac{\pi}{P_3} |x_i - x_j|\right]\right). \end{aligned}$$

The periods  $P_1$  and  $P_2$  were set to match the data periodicity found in Task 1, with

$$\begin{aligned} P_2 &= 0.229 \text{ day}^{-1}, \text{ and} \\ P_3 &= 0.421 \text{ day}^{-1}, \end{aligned}$$

while the scale of the correlations were initially set to

$$\Gamma_2 = \Gamma_3 = 1.$$

The Gaussian process was computed before optimising the parameters. Figure 7(a) shows predictions for the future brightness of this star for two months after the last data point. Figure 7(b) shows the same prediction after optimising over the parameters  $P_2$ ,  $P_3$ ,  $\Gamma_2$  and  $\Gamma_3$ . Notably,  $P_2$  and  $P_3$  were *not* kept constant during optimisation, as the Lomb-Scargle periodogram (which was used to calculate their initial values) did not take into account the  $y$ -error. As can be seen in from Figures 7(a) and 7(b), the optimisation did not significantly change the predictions, however the predictions look sensible nonetheless. The code for Question 2 is shown in Appendix B.

### 3 Question 3

#### 3.1 Task 1

In this task, we build and train an autoencoder from images of hand written digits. Our autoencoder was a fully connected dense neural network with 5 layers, starting with the  $28 \times 28$  input, 3 hidden layers, and a  $28 \times 28$  output. The middle layer was an 8 dimensional latent space. Specifically our autoencoder model is:

$$\mathbf{x} = \text{AE}(\mathbf{x}),$$


where,

$$\text{AE}(\mathbf{x}) = \text{OUT}(H_2(L(H_1(\mathbf{x})))),$$


and,

$$\begin{aligned} H_1(\mathbf{x}) &= \text{ReLU}(W_{H_1}\mathbf{x} + \mathbf{b}_{H_1}), \\ L(\mathbf{x}) &= \text{ReLU}(W_L\mathbf{x} + \mathbf{b}_L), \\ H_2(\mathbf{x}) &= \text{Sigmoid}(W_{H_2}\mathbf{x} + \mathbf{b}_{H_2}), \\ \text{OUT}(\mathbf{x}) &= \text{Sigmoid}(W_O\mathbf{x} + \mathbf{b}_O). \end{aligned}$$

Weight matrix  $W_{H_1}$  has size  $(32 \times 784)$ ,  $W_L$  has size  $(8 \times 32)$ ,  $W_{H_2}$  has size  $(32 \times 8)$ , and  $W_O$  has size  $(784 \times 32)$ . Similarly, bias vector  $\mathbf{b}_{H_1}$  is 32 dimensional,  $\mathbf{b}_L$  is 8 dimensional,  $\mathbf{b}_{H_2}$  is 32 dimensional,  $\mathbf{b}_O$  is 784 dimensional. Training the network involves learning the elements of the aforementioned matrices and vectors, such that the output,  $\text{AE}(\mathbf{x})$ , most closely matches  $\mathbf{x}$ . The autoencoder was trained for 50 epochs with a batch size of 256. Figure 8 shows some sample predictions of the autoencoder after training. The latent space (output of  $L(H_1(\mathbf{x}))$ ) is shown in Figure 9.



(a)



(b)

Figure 7: (a) The Gaussian process predictions before optimisation. (b) The Gaussian process predictions after optimisation.




Figure 8: Sample inputs (top) and autoencoder predictions (bottom).

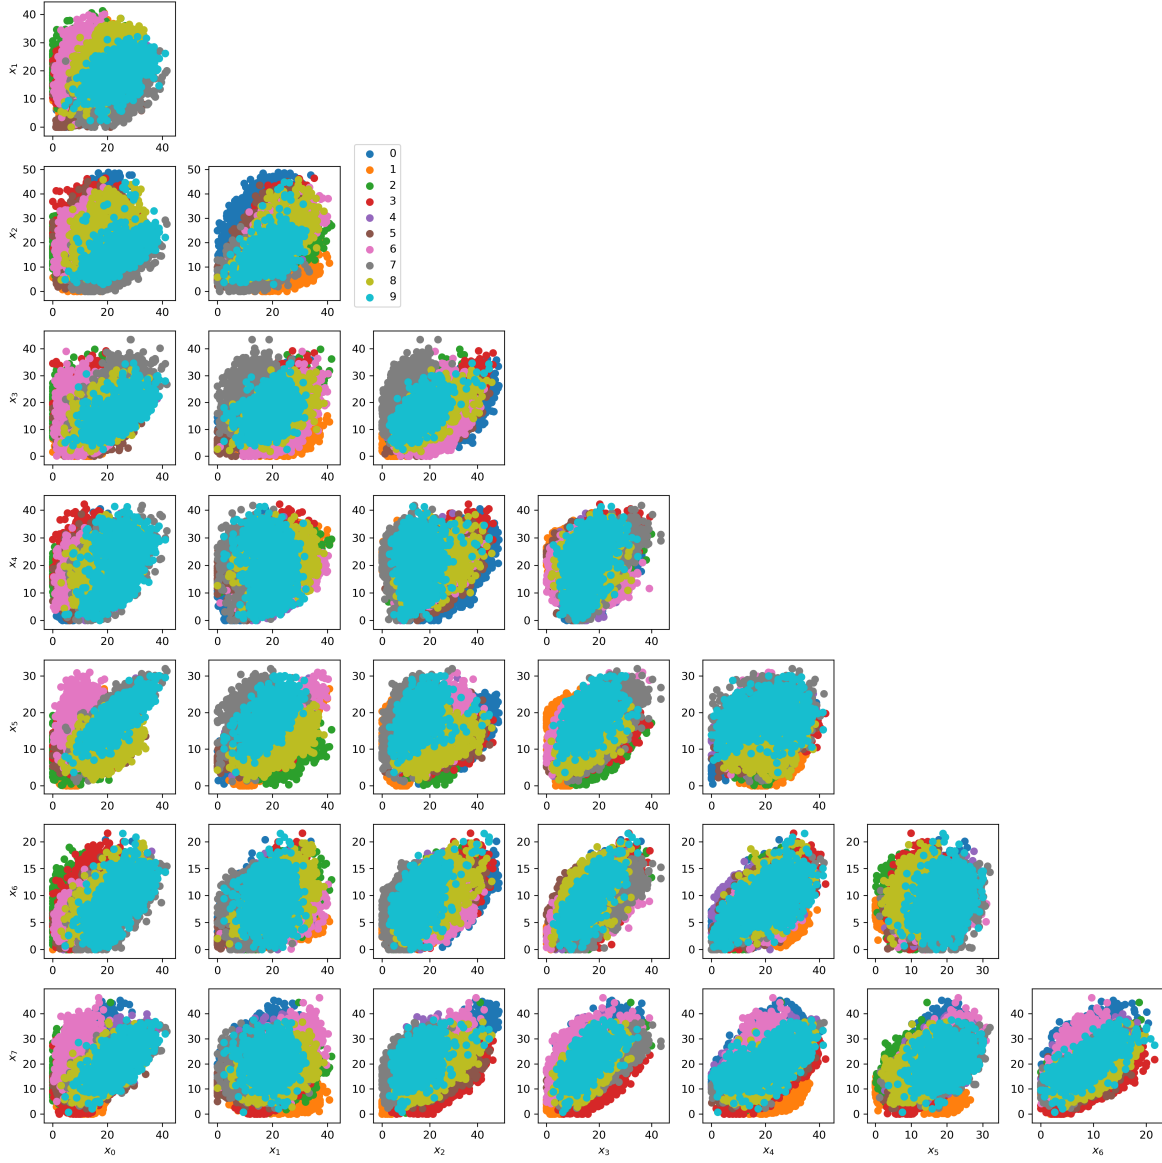


Figure 9: The autoencoder latent space.






Figure 10: The autoencoder predictions at one slice of latent space.

### 3.2 Task 2

Figure 9 shows the autoencoder predictions across two dimensions of latent space. Since the latent space is 8 dimensional, the remaining 6 dimensions were kept constant. Specifically, we had

$$\begin{aligned}
 x_1 &= 20, \\
 x_2 &= 20, \\
 x_3 &= 10, \\
 x_4 &= 20, \\
 x_6 &= 20, \text{ and} \\
 x_7 &= 20.
 \end{aligned}$$

while  $x_0$  was plotted against  $x_5$ . These values were chosen so that the predictions were from latent space values near the data shown in Figure 10. The code for Question 3 is shown in Appendix C.

## A Question 1

```
1 # imports
2 import csv
3 import urllib.request as request
4 import numpy as np
5 import matplotlib.pyplot as plt
6 from matplotlib.ticker import MaxNLocator
7 import pystan as stan
8
9
10 # get data
11 url = "http://astrowizici.st/teaching/phs5000/decay.csv"
12 response = request.urlopen(url)
13 lines = [l.decode('utf-8') for l in response.readlines()]
14 cr = csv.reader(lines)
15 cr_data = list(cr)
16 data = np.array(cr_data[1:])
17 t, N, N_err, d_name = data.T
18 # time in seconds
19 t = t.astype(np.float)
20 # grams of radioactive material measured
21 N = N.astype(np.float)
22 # uncertainty in grams measured
23 N_err = N_err.astype(np.float)
24 # Detector name
25 d_name = d_name.astype(str)
26
27 # define useful quantities
28 days_to_s = 60 * 60 * 24
29 s_to_days = 1/days_to_s
30 # known parameters:
31 N_widgets = 100
32 N_initial_max = 20
33 manufacturing_time_span = 35 * days_to_s
34 t_delay = 14 * days_to_s
35 measurement_time_span = 90 * days_to_s
36 N_observations = 1
37 # time in days
38 t_days = t * s_to_days
39 # minimum measured N
40 min_N = np.min(N)
41 # Minimum time span:
42 min_dt = t_delay
43 # maximum possible alpha value
44 max_alpha = -1/(min_dt) * np.log(min_N/N_initial_max)
45
46
47 # plot figure
48 fig, ax = plt.subplots(figsize=(4, 4))
49 ax.scatter(t_days, N, c="k", s=10)
50 ax.errorbar(t_days, N, yerr=N_err, fmt="o", lw=1, c="k")
51 ax.axhline(0, linestyle='--', c='gray')
52 ax.set_xlabel(r"Time [days since manufacturing]")
53 ax.set_ylabel(r"Measured Radioactive material [grams]")
54 ax.xaxis.set_major_locator(MaxNLocator(6))
55 ax.yaxis.set_major_locator(MaxNLocator(7))
56 fig.tight_layout()
57 plt.savefig('q1_data.png', dpi=300)
58
59
60 # pystan utils
61 def sampling_kwds(**kwargs):
62     r"""
63     Prepare a dictionary that can be passed to Stan at the sampling stage.
64     Basically this just prepares the initial positions so that they match the
65     number of chains.
66     """
67
68     kwds = dict(chains=4)
69     kwds.update(kwargs)
```

```

70
71     if "init" in kwds:
72         kwds["init"] = [kwds["init"]] * kwds["chains"]
73
74     return kwds
75
76
77 # define model
78 model_filename = "model1.stan"
79 model_str = """
80 data {
81     int<lower=1> N_widgets; // number of widgets
82
83     // Time of measurement.
84     vector[N_widgets] t_measured;
85
86     // Amount of material measured is uncertain.
87     vector[N_widgets] N_measured;
88     vector[N_widgets] sigma_N_measured;
89
90     // Maximum amount of initial material.
91     real N_initial_max;
92
93     // last possible time of manufacture.
94     real t_initial_max;
95
96     // max value of alpha
97     real alpha_max;
98 }
99
100 parameters {
101     // Time of manufacture.
102     vector<lower=0, upper=t_initial_max>[N_widgets] t_initial;
103
104     // The decay rate parameter.
105     real<lower=0, upper=alpha_max> alpha;
106
107     // The amount of initial material is not known.
108     vector<lower=0, upper=N_initial_max>[N_widgets] N_initial;
109 }
110
111 model {
112     for (i in 1:N_widgets) {
113         N_measured[i] ~ normal(
114             N_initial[i] * exp(-alpha * (t_measured[i] - t_initial[i])),
115             sigma_N_measured[i]
116         );
117     }
118 }
119 """
120 # make model
121 model = stan.StanModel(model_code=model_str)
122
123 # Data.
124 data_dict = dict(
125     N_widgets=N_widgets,
126     t_measured=t,
127     N_measured=N,
128     sigma_N_measured=N_err,
129     N_initial_max=N_initial_max,
130     t_initial_max=manufacturing_time_span,
131     alpha_max=max_alpha,
132 )
133
134 # initial guess
135 alpha_guess = max_alpha/2
136 t_init_guess = np.full(N_widgets, manufacturing_time_span/2)
137
138 init_dict = dict(
139     t_initial=t_init_guess,
140     alpha=alpha_guess,

```

```

141 )
142
143 # Run optimisation.
144 opt_stan = model.optimizing(
145     data=data_dict,
146     init=init_dict
147 )
148
149 # Run sampling.
150 samples = model.sampling(**sampling_kwds(
151     chains=2,
152     iter=2000,
153     data=data_dict,
154     init=opt_stan
155 ))
156
157
158 # plot initial N
159 fig = samples.traceplot(("N_initial", ))
160 fig.set_size_inches(10, 5)
161 plt.savefig('q11_N_init.png', dpi=300)
162
163 # plot initial t
164 fig = samples.traceplot(("t_initial", ))
165 fig.set_size_inches(10, 5)
166 plt.savefig('q11_t_init.png', dpi=300)
167
168
169 # get alpha chain and remove burn in
170 alpha_chain = samples["alpha"][1000:]
171 # calculate mean and std
172 alpha_mean = np.mean(alpha_chain)
173 alpha_err = np.std(alpha_chain)
174
175 # plot
176 fig, axs = plt.subplots(1, 2, figsize=(12, 5))
177
178 ax = axs[0]
179 ax.plot(alpha_chain, c='k')
180 ax.set_xlabel("Step Number")
181 ax.set_ylabel("Alpha")
182
183 ax = axs[1]
184 ax.hist(alpha_chain, bins=30, color='gray', label='Posterior Distribution')
185 ax.axvline(alpha_mean, c='k', linestyle='--', label='Mean')
186 ax.set_ylabel("Posterior Probability")
187 ax.set_xlabel("Alpha")
188 ax.legend()
189 plt.savefig('q11_alpha.png', dpi=300)
190
191
192 # print mean and std
193 print(f'Alpha: {alpha_mean:.2}, Standard Deviation: {alpha_err:.0}')
194
195
196 # Part 2
197
198
199 # make model
200 model_filename = "model2.stan"
201 model_str = """
202 data {
203     int<lower=1> N_widgets; // number of widgets
204
205     // Time of measurement.
206     vector[N_widgets] t_measured;
207
208     // Amount of material measured is uncertain.
209     vector[N_widgets] N_measured;
210     vector[N_widgets] sigma_N_measured;
211

```

```

212 // Maximum amount of initial material.
213 real N_initial_max;
214
215 // last possible time of manufacture.
216 real t_initial_max;
217
218 // max value of alpha
219 real alpha_max;
220
221 //detector
222 int k[N_widgets];
223 }
224
225 parameters {
226 // Time of manufacture.
227 vector<lower=0, upper=t_initial_max>[N_widgets] t_initial;
228
229 // The decay rate parameter.
230 real<lower=0, upper=alpha_max> alpha;
231
232 // The detector bias.
233 real bias;
234
235 // The amount of initial material is not known.
236 vector<lower=0, upper=N_initial_max>[N_widgets] N_initial;
237
238 // multinomial
239 simplex[3] theta;
240 }
241
242 model {
243   for (i in 1:N_widgets) {
244     N_measured[i] ~ normal(
245       N_initial[i] * exp(-alpha * (t_measured[i] - t_initial[i]))
246       + theta[k[i]] * bias,
247       sigma_N_measured[i]
248     );
249   }
250 }
251 """
252 model = stan.StanModel(model_code=model_str)
253
254
255 # convert detector names to numbers
256 name_to_k = dict(
257   A=1,
258   B=2,
259   C=3,
260 )
261 k = np.vectorize(name_to_k.get)(d_name)
262
263 # Data.
264 data_dict = dict(
265   N_widgets=N_widgets,
266   t_measured=t,
267   N_measured=N,
268   sigma_N_measured=N_err,
269   N_initial_max=N_initial_max,
270   t_initial_max=manufacturing_time_span,
271   alpha_max=max_alpha,
272   k=k,
273 )
274
275 # initial guess
276 alpha_guess = max_alpha/2
277 t_init_guess = np.full(N_widgets, manufacturing_time_span/2)
278 theta_guess = [0.33, 0.33, 0.34]
279 bias_guess = 0.
280
281 init_dict = dict(
282   t_initial=t_init_guess,

```

```

283     alpha=alpha_guess,
284     theta=theta_guess,
285     bias=bias_guess,
286 )
287
288 # Run optimisation.
289 opt_stan = model.optimizing(
290     data=data_dict,
291     init=init_dict
292 )
293
294 # Run sampling.
295 samples = model.sampling(**sampling_kwds(
296     chains=2,
297     iter=2000,
298     data=data_dict,
299     init=opt_stan,
300     control=dict(max_treedepth=30),
301 ))
302
303
304 # plot initial N
305 fig = samples.traceplot(("N_initial", ))
306 fig.set_size_inches(10, 5)
307 plt.savefig('q12_N_init.png', dpi=300)
308
309 # plot initial t
310 fig = samples.traceplot(("t_initial", ))
311 fig.set_size_inches(10, 5)
312 plt.savefig('q12_t_init.png', dpi=300)
313
314
315 # get theta chains and remove burn in
316 theta_chains = samples["theta"][1000:].T
317
318
319 # plot theta chains
320 fig, axs = plt.subplots(1, 2, figsize=(12, 5))
321
322 ax = axs[0]
323 for i in range(3):
324     ax.plot(theta_chains[i])
325 ax.set_xlabel("Step Number")
326 ax.set_ylabel(r"$\theta_k$")
327
328 ax = axs[1]
329 for i in range(3):
330     ax.hist(theta_chains[i], bins=30, label=rf'$\theta_{i+1}$')
331 ax.legend()
332 ax.set_ylabel("Posterior Probability")
333 ax.set_xlabel(r"$\theta_k$")
334 plt.savefig('q12_theta.png', dpi=300)
335
336
337 # get bias chain and remove burn in
338 bias_chain = samples["bias"][1000:]
339 # calculate mean and std
340 bias_mean = np.mean(bias_chain)
341 bias_err = np.std(bias_chain)
342
343 # plot
344 fig, axs = plt.subplots(1, 2, figsize=(12, 5))
345
346 ax = axs[0]
347 ax.plot(bias_chain, c='k')
348 ax.set_xlabel("Step Number")
349 ax.set_ylabel("Bias")
350
351 ax = axs[1]
352 ax.hist(bias_chain, bins=30, color='gray', label='Posterior Distribution')
353 ax.axvline(bias_mean, c='k', linestyle='--', label='Mean')

```

```

354 ax.set_ylabel("Posterior Probability")
355 ax.set_xlabel("Bias")
356 ax.legend()
357 plt.savefig('q12_bias.png', dpi=300)
358
359
360 # print mean and std
361 print(f'Bias: {bias_mean:.2}, Standard Deviation: {bias_err:.0}')
362
363
364 # get alpha chain and remove burn in
365 alpha_chain = samples["alpha"][1000:]
366 # calculate mean and std
367 alpha_mean = np.mean(alpha_chain)
368 alpha_err = np.std(alpha_chain)
369
370 # plot
371 fig, axs = plt.subplots(1, 2, figsize=(12, 5))
372
373 ax = axs[0]
374 ax.plot(alpha_chain, c='k')
375 ax.set_xlabel("Step Number")
376 ax.set_ylabel("Alpha")
377
378 ax = axs[1]
379 ax.hist(alpha_chain, bins=30, color='gray', label='Posterior Distribution')
380 ax.axvline(alpha_mean, c='k', linestyle='--', label='Mean')
381 ax.set_ylabel("Posterior Probability")
382 ax.set_xlabel("Alpha")
383 ax.legend()
384 plt.savefig('q12_alpha.png', dpi=300)
385
386
387 # print mean and std
388 print(f'Alpha: {alpha_mean:.2}, Standard Deviation: {alpha_err:.0}')

```

## B Question 2

```

1 import scipy.optimize as op
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from matplotlib.ticker import MaxNLocator
5 import pickle
6 from scipy.signal import lombscargle
7 from george import kernels
8 import george
9
10 # load data
11 with open("assignment2_gp.pkl", "rb") as fp:
12     data = pickle.load(fp)
13     t = data['t']
14     y = data['y']
15     y_err = data['yerr']
16
17 # plot the data
18 fig, ax = plt.subplots(figsize=(8, 5))
19 ax.scatter(t, y, c="k", s=2)
20 ax.set_xlabel(r"$t$")
21 ax.set_ylabel(r"$y$")
22 ax.xaxis.set_major_locator(MaxNLocator(6))
23 ax.yaxis.set_major_locator(MaxNLocator(7))
24 fig.tight_layout()
25 plt.savefig('q2_data.png', dpi=300)
26
27
28 # get periodogram
29 f = np.linspace(0.001, 2, 1000)
30 pgram = lombscargle(t, y, f)
31
32

```

```

33 # find index of peaks
34 peak_i = np.where(
35     np.r_[True, pgram[1:] > pgram[:-1]] & np.r_[pgram[:-1] > pgram[1:], True]
36 )[0]
37 # power of peaks
38 peak_p = pgram[peak_i]
39
40 # exclude peaks with <100 power
41 peak_i = peak_i[np.where(peak_p > 100)]
42 peak_p = pgram[peak_i]
43 # frequency of peaks
44 peak_f = f[peak_i]
45
46 # plot
47 plt.figure(figsize=(8, 5))
48 plt.plot(f, pgram, c='k')
49 plt.xlabel("Frequency")
50 plt.ylabel("Power")
51 plt.vlines(peak_f, 0, peak_p, linestyle='--', colors='gray')
52 plt.savefig('periodogram.png', dpi=300)
53
54 # get periods (for later)
55 periods = 2*np.pi / peak_f
56 print('Frequency: \t Power: \t Periods:')
57 for freq, power, T in zip(peak_f, peak_p, periods):
58     print(f'{freq:.3}\t \t {int(power)} \t \t {int(T)}')
59
60
61 # Part 2
62
63
64 # define our kernal
65 k1 = kernels.ExpSine2Kernel(gamma=1, log_period=np.log(periods[1]))
66 k2 = kernels.ExpSine2Kernel(gamma=1, log_period=np.log(periods[3]))
67
68 kernel = k1 + k2
69 kernel *= 1 * kernels.ExpSquaredKernel(1000)
70
71 # compute the gaussian process
72 gp = george.GP(kernel)
73 gp.compute(t, y_err)
74
75 t_min = np.min(t)
76 t_max = np.max(t)
77 t_pred = np.linspace(t_min, t_max + 60, 500)
78 pred, pred_var = gp.predict(y, t_pred, return_var=True)
79
80
81 plt.figure(figsize=(8, 5))
82 plt.fill_between(t_pred, pred - np.sqrt(pred_var), pred + np.sqrt(pred_var),
83                 color="k", alpha=0.2)
84 plt.plot(t_pred, pred, "k", lw=1.5, alpha=0.5)
85 plt.scatter(t, y, c="k", s=2)
86 plt.xlim(t_min, t_max + 60)
87 plt.xlabel("Day")
88 plt.ylabel("Brightness")
89 plt.savefig('gp1.png', dpi=300)
90
91 print(f"Initial ln-likelihood: {gp.log_likelihood(y):.2f}")
92
93
94 # Define the objective function (negative log-likelihood in this case).
95 def nll(p):
96     gp.set_parameter_vector(p)
97     ll = gp.log_likelihood(y, quiet=True)
98     return -ll if np.isfinite(ll) else 1e25
99
100
101 # And the gradient of the objective function.
102 def grad_nll(p):
103     gp.set_parameter_vector(p)

```



```

104     return -gp.grad_log_likelihood(y, quiet=True)
105
106
107 # Run the optimization routine.
108 p0 = gp.get_parameter_vector()
109 results = op.minimize(nll, p0, jac=grad_nll, method="L-BFGS-B")
110
111 # Update the kernel and print the final log-likelihood.
112 gp.set_parameter_vector(results.x)
113 print(f"Final ln-likelihood: {gp.log_likelihood(y):.2f}")
114
115 t_min = np.min(t)
116 t_max = np.max(t)
117 t_pred = np.linspace(t_min, t_max + 60, 500)
118 mu, var = gp.predict(y, t_pred, return_var=True)
119
120
121 # plot
122 plt.figure(figsize=(8, 5))
123 plt.fill_between(t_pred, pred - np.sqrt(pred_var), pred + np.sqrt(pred_var),
124                 color="k", alpha=0.2)
125 plt.plot(t_pred, pred, "k", lw=1.5, alpha=0.5)
126 plt.scatter(t, y, c="k", s=2)
127 plt.xlim(t_min, t_max + 60)
128 plt.xlabel("Day")
129 plt.ylabel("Brightness")
130 plt.savefig("gp2.png", dpi=300)

```

## C Question 3

```

1 import gzip
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from tensorflow import keras
5 from tensorflow.keras import layers
6 import tensorflow as tf
7
8
9 tf.random.set_seed(1)
10
11
12 # training/test set size
13 n_train = 60000
14 n_test = 10000
15
16
17 # get data function
18 def get_data(fname):
19     image_size = 28
20     f = gzip.open(fname, 'r')
21     f.read(16)
22     data = f.read()
23     data = np.frombuffer(data, dtype=np.uint8).astype(np.float32)
24     data = data.reshape(-1, image_size, image_size)
25     f.close()
26     return data
27
28
29 # get labels function
30 def get_labels(fname):
31     f = gzip.open(fname, 'r')
32     f.read(8)
33     data = f.read()
34     data = np.frombuffer(data, dtype=np.uint8)
35     f.close()
36     return data
37
38
39 # get data
40 train_data = get_data('train-images-idx3-ubyte.gz')

```

```

41 test_data = get_data('t10k-images-idx3-ubyte.gz')
42 labels = get_labels('t10k-labels-idx1-ubyte.gz')
43
44
45 # specify network
46 input_img = keras.Input(shape=(28*28,))
47
48 h_size = 32
49 latent_size = 8
50
51 h1 = layers.Dense(h_size, activation='relu')(input_img)
52 encoded = layers.Dense(latent_size, activation='relu')(h1)
53
54 h2 = layers.Dense(h_size, activation='relu')(encoded)
55 decoded = layers.Dense(28*28, activation='sigmoid')(h2)
56
57 autoencoder = keras.Model(input_img, decoded)
58
59 # just the encoder part
60 encoder = keras.Model(input_img, encoded)
61
62 # just the decoder part
63 encoded_input = keras.Input(shape=(latent_size,))
64 h2_layer = autoencoder.layers[-2](encoded_input)
65 decoder_layer = autoencoder.layers[-1](h2_layer)
66 decoder = keras.Model(encoded_input, decoder_layer)
67
68
69 # compile network
70 autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
71
72 # normalise and flatten data
73 train_data = (train_data/255).reshape((-1, 28*28))
74 test_data = (test_data/255).reshape((-1, 28*28))
75
76
77 # training
78 autoencoder.fit(train_data, train_data,
79                 epochs=50,
80                 batch_size=256,
81                 shuffle=True,
82                 validation_data=(test_data, test_data))
83
84
85 # make predictions
86 encoded_imgs = encoder.predict(test_data)
87 decoded_imgs = decoder.predict(encoded_imgs)
88
89 n = 10
90 plt.figure(figsize=(20, 4))
91 for i in range(n):
92     # Display original
93     ax = plt.subplot(2, n, i + 1)
94     plt.imshow(test_data[i].reshape(28, 28))
95     plt.gray()
96     ax.get_xaxis().set_visible(False)
97     ax.get_yaxis().set_visible(False)
98
99     # Display reconstruction
100    ax = plt.subplot(2, n, i + 1 + n)
101    plt.imshow(decoded_imgs[i].reshape(28, 28))
102    plt.gray()
103    ax.get_xaxis().set_visible(False)
104    ax.get_yaxis().set_visible(False)
105
106 plt.tight_layout()
107 plt.savefig('ae_predictions.png', dpi=300)
108
109
110 # show latent space
111 x_arr = np.arange(8)

```

```

112 y_arr = np.arange(2, 4)
113 mesh = np.array(np.meshgrid(x_arr, y_arr))
114 combinations = mesh.T.reshape(-1, 2)
115
116 arrs = encoded_imgs.T
117 arrs = encoded_imgs.T
118 fig, axs = plt.subplots(7, 7, figsize=(15, 15))
119 label = False
120 for i in range(8):
121     for j in range(7):
122         ax = axs[i-1, j]
123         if (j >= i):
124             ax.axis('off')
125             continue
126         ax.axis('on')
127         x = arrs[j]
128         y = arrs[i]
129         for n in range(10):
130             k = np.where(labels == n)
131             if not label:
132                 ax.scatter(x[k], y[k], label=n)
133             else:
134                 ax.scatter(x[k], y[k])
135         label = True
136         if i == 7:
137             ax.set_xlabel(rf"$x_{j}$")
138         if j == 0:
139             ax.set_ylabel(rf"$x_{i}$")
140 fig.legend(loc='center left', bbox_to_anchor=(0.3, 0.8))
141 plt.tight_layout()
142 plt.savefig('latent_space.png', dpi=300)
143
144
145 # Save models
146 autoencoder.save('autoencoder')
147 encoder.save('encoder')
148 decoder.save('decoder')
149
150
151 # part 2
152
153 # number of images = nx * ny
154 nx = 15
155 ny = 15
156 # how much of latent space we will plot
157 x_min = 0
158 x_max = 40
159 y_min = 0
160 y_max = 30
161
162 xs = np.linspace(x_min, x_max, nx)
163 ys = np.linspace(y_min, y_max, ny)
164
165 # will plot how these change
166 x0, x5 = np.meshgrid(xs, ys)
167 shape = x0.shape
168
169 # get constant values for each other node
170 x1 = np.full(shape, 20)
171 x2 = np.full(shape, 20)
172 x3 = np.full(shape, 10)
173 x4 = np.full(shape, 20)
174 x6 = np.full(shape, 10)
175 x7 = np.full(shape, 20)
176
177 full_grid = np.stack((x0, x1, x2, x3, x4, x5, x6, x7)).T
178
179 # reshape into big combined picture
180 full_grid = full_grid.reshape(nx*ny, 8)
181 decoded_grid = decoder.predict(full_grid)
182 decoded_grid = decoded_grid.reshape((nx, ny, 28, 28))

```

```
183 decoded_grid = decoded_grid.swapaxes(1, 2)
184 decoded_grid = decoded_grid.reshape(28*nx, 28*ny)
185
186 fig, ax = plt.subplots(figsize=(10, 10))
187 extent = [x_min, x_max, y_min, y_max]
188 ax.imshow(decoded_grid, extent=extent)
189 ax.set_xlabel(r"$x_0$")
190 ax.set_ylabel(r"$x_5$")
191 plt.tight_layout()
192 plt.savefig('image_grid.png', dpi=300)
```